

Title

EN source - machine-translatable excerpt

The question of local inference has grown less rhetorical and more practical in the last
There are three reasons an institution keeps its data on a single machine. The first is confidentiality: a manuscript must never cross a public network. The second is latency: a translator will not watch a spinner for twenty seconds. The third is cost, which
Consider the tokenization problem. A model that counts words as we do will regularly misjudge the length of a paragraph, sometimes by a factor of two. This matters more than it appears, because budgeting is how we avoid the failure mode in which the
Footnote: see also the discussion of context windows in chapter four, where the author argues that a planner must reserve output headroom rather than assume the model
The evidence, such as it is, points toward hybrid strategies. A small local model can handle the mechanical work - segmentation, normalization, extraction - while a larger model is reserved for the judgment that still resists automation. This division of labour mirrors, in its own way, the older practice of the translator's assistant reading the

Notes

See also the discussion of context windows in chapter four.